

Faraz Mazhar

Managing Consultant — Data Engineering
Lahore, Pakistan
farazmazhar.com

csfarazmazhar@gmail.com
linkedin.com/in/faraz-mazhar
github.com/farazmazhar

PROFESSIONAL SUMMARY

Data engineer with 8+ years building scalable data platforms on AWS and Azure. Architected and led development of a modular data framework powering enterprise analytical and ML pipelines. Owns the full lifecycle from data modeling and lakehouse design through API integration, orchestration, and production deployment.

TECHNICAL SKILLS

Cloud	AWS (S3, Lambda, Glue, EMR, API Gateway, Step Functions, CloudFormation, DynamoDB, RDS, Athena, CloudWatch, IAM, KMS, SNS, DMS), Azure (Data Factory, Databricks, Functions, EventHub, App Service, SQL DB, Active Directory)
Big Data	Apache Spark, Delta Lake, Databricks, Apache Airflow, Kafka, Hadoop (Hive, HDFS, Yarn), Trino/Presto, Sqoop
Languages	Python (Django, Flask, Pandas, PySpark), SQL, Bash, PowerShell, C/C++
Databases	SQL Server, PostgreSQL (RDS, Redshift), MongoDB, DynamoDB
APIs & IaC	REST, GraphQL (Strawberry-Django), API Gateway, Kinesis, CloudFormation, Git, Linux

EXPERIENCE

Systems Limited

Mar 2021 — Present

Managing Consultant

Python · Pandas · PySpark · Delta Lake · Azure · Databricks · Airflow · Django

- Led solution architecture and development of a modular Data Framework (Python, Pandas, PySpark, Delta Lake) for ingestion, curation, and analytics — enabling data pipelines without deployment cycles.
- Owned end-to-end solution architecture across the platform — designed features spanning data ingestion, migration, reconciliation, policy frameworks, alerting, automated reporting, visual analytics, and platform-wide enhancements and optimizations.
- Designed a DSL-driven Canonical Model and Semantic Model for schema-agnostic data curation across arbitrary inputs and outputs.
- Data-modeled the metastore, Delta Lake, and SQL Server datastore; optimized for query performance and multi-tenancy.
- Architected REST and GraphQL APIs for Integration and Semantic services, enabling external systems to submit workloads and consume metadata dynamically.
- Orchestrated multi-stage pipelines with Databricks Workflows, Apache Airflow, and Azure Data Factory; built a dialect-agnostic SQL query builder.

Northbay Solutions

Aug 2018 — Mar 2021

Senior Software Engineer

Python · PySpark · AWS Glue · Lambda · Step Functions · CloudFormation · Flask

- Migrated legacy SAS code to PySpark and built reusable packages for AWS Glue ETL jobs; developed a Data Lake framework for large-scale ingestion and processing.
- Built event-driven ETL triggers using Lambda and CloudWatch on S3 events, and API Gateway endpoints for Kinesis Data Streams with DynamoDB tooling.
- Designed dynamically parallel Step Functions State Machines and built an automated cloud migration framework with CloudFormation IaC.
- Developed REST APIs with Flask and a customizable XML parser for document processing.

Mentor Graphics

2017

Intern — Quality Assurance

C/C++ · Python · Bash · Nucleus RTOS

- Performed QA and test automation for Nucleus RTOS; wrote test scripts in Python and Bash.

EDUCATION

BS (Hons.) Computer Science

Punjab University College of Information Technology (PUCIT)
2014 — 2018

CERTIFICATIONS

Microsoft Azure Fundamentals (AZ-900)

2026

AWS Certified Developer — Associate

2018 — 2021